

R-on-Rails: ship the ledger on Rails, keep the science in R

A portfolio stress tester that puts SQLite, Hotwire, and GNU R on one desk

Rodrigo Botafogo

2026

Contents

1	Two audiences, one gap	2
2	What the Ledger is	2
3	What's in Rails (and the database)	3
4	R calculators that read like Ruby	3
5	How data moves (Arrow, honestly)	4
5.1	Example: copy into R	4
5.2	Example: IPC file handoff	5
6	The page with the charts	6
6.1	Side-by-side KPIs	6
6.2	Density and tail risk	6
6.3	Thirty-day Monte Carlo cone	6
6.4	Rolling historical VaR and breaches	6
7	Why this is a production-shaped story	6
8	Try it	9
9	Honest limits	9



1 Two audiences, one gap

If you live in **R**, you already have the hard part: VaR, density estimates, Monte Carlo, Bioconductor, the whole CRAN catalog. What is still painful is turning a notebook into a **product**—logins, a database, background jobs, a page that updates when the job finishes, and a deploy story that is not a pile of glue scripts.

If you live in **Ruby on Rails**, you already have the product shell. What is still painful is serious statistics without standing up a Python microservice fleet, learning another ORM, and ferrying CSV files between processes.

R-on-Rails is the name we give to a simple division of labor:

- **Rails** owns HTTP, Active Record, jobs, Hotwire, and the UX.
- **GNU R** owns the math and the statistical vocabulary.
- **Galaaz** is the bridge that makes R feel like a Ruby DSL—and moves bulky tables with Apache Arrow when you need more than a few numbers.

This post walks through a real demo app—the **R-on-Rails Ledger**— so you can see the database, the calculators, and the charts on one page.

2 What the Ledger is

The Ledger is a **family-office style portfolio stress tester** on **Rails 8**, **CRuby**, **SQLite**, **Solid Queue** / **Solid Cable**, and **Hotwire**. It is a sibling app to the Galaaz gem (not packaged inside the gem). One click runs two engines on the same portfolio return panel:

- **Engine A — historical:** empirical VaR / expected shortfall / return density from history (`R.quantile`, `R.mean`, `R.density`).
- **Engine B — Monte Carlo:** forward GBM paths calibrated to that history, then forward VaR-style KPIs and a 30-day “cone.”

Rails stores the book and the job results; R does the science; the browser gets Plotly charts from JSON that R (and a little Ruby) produced.



3 What's in Rails (and the database)

The “ledger” here is a **portfolio price book**, not double-entry accounting:

Table	Role
portfolios	Named book (name, total_value)
assets	Tickers + weights
historical_prices	Bars: adjusted_close, daily_return, ...
stress_tests	Job status, scalar KPIs, chart_payload JSON
stress_test_runs	Per-engine (historical / monte_carlo) payloads

Seed data is **synthetic GBM** (honest demo data—not a live feed). `SEED_PROFILE=fast` is laptop-friendly (~10 assets x 2k bars); `SEED_PROFILE=wow` is the “about a million bars” line.

The click path is ordinary Rails:

```
POST /portfolios/:id/stress_tests
-> StressTestsController#create
-> StressTestJob.perform_later(...)
-> Risk::DualOrchestrator
    |- Risk::HistoricalEngine
    |- Risk::MonteCarloEngine
-> Turbo::StreamsChannel broadcast
-> results partial + Plotly redraw
```

No React SPA. No Redis required for the demo. Development often runs the job `:async` in-process so one `bin/dev` is enough; Solid Queue stays in the Gemfile for production-shaped runs.

4 R calculators that read like Ruby

After returns land in R (via Arrow—see below), the historical engine exposes small methods that call R the way Galaaz intends: **named functions and vectors**, not a wall of quoted R source for every KPI.

From the Ledger’s `Risk::HistoricalEngine` (abbreviated):

```
def historical_var(returns_vec, probs:)
  RValues.scalar_f(
    R.as_numeric(
      R.quantile(returns_vec,
                  probs: probs,
                  names: false,
                  type: 7)
    )
  )
end

def expected_shortfall(returns_vec, var_level)
```

```
thr = var_level.is_a?(Numeric) ?  
  var_level : RValues.scalar_f(var_level)  
RValues.scalar_f(  
  R.as_numeric(  
    R.mean(returns_vec[returns_vec <= thr])  
  )  
)  
end  
  
def return_density(returns_vec, n: 128)  
  ArrowHandoff.density_xy_for_plotly(returns_vec, n: n)  
end
```

That is the pitch for Ruby developers: `R.quantile` is a method call. For R developers: the quantile is still **R**'s quantile—same `type = 7` story you know—just invoked from the app that owns the database.

A tiny live taste of the same idea (outside Rails):

```
x = R.c(-0.02, -0.01, 0.0, 0.01, 0.015, -0.005)  
q05 = R.as_numeric(  
  R.quantile(x, probs: 0.05, names: false, type: 7))  
puts "5% quantile: #{q05 >> 0}"
```

```
## 5% quantile: -0.017499999999999998
```

Engine B calibrates μ/σ in R, simulates GBM paths, and returns structured fields (`sim_var_30d`, sample paths, terminal density) that Rails stores on each run's payload JSON.

5 How data moves (Arrow, honestly)

Crossing a process boundary means you must be honest about **copies**. Galaaz can move columnar data with Apache Arrow in three ways:

1. **Copy into R** — Ruby builds batches; R materialises an Arrow table / proxy. Simple and available today.
2. **IPC file handoff** — Ruby writes an Arrow IPC file (often under `/dev/shm`); R opens it by **path**. Only the path crosses NewBridge—not megabytes of `MsgPack`. This is what the Ledger prefers for the return panel.
3. **Shared-memory zero-copy** — both processes attach to one live segment. **Not implemented yet**; we do not claim it.

5.1 Example: copy into R

Good when the table is modest or you do not have a Ruby Arrow writer installed. One call builds the R-side table; then you Remote-Control it.



```
arrow_ok = R::Support.eval(
  "requireNamespace('arrow', quietly=TRUE)" == true
unless arrow_ok
  puts '(Skip: need R package arrow.)'
else
  batches = [
    [{ daily_return: -0.01 }, { daily_return: 0.02 }],
    [{ daily_return: 0.005 }, { daily_return: -0.003 }]
  ]
  tbl = R::Arrow.from_ruby_batches(batches)
  puts "R class: #{tbl.rclass}"
  vec = R.as__numeric(
    R.dplyr__collect(tbl)[['daily_return']]
  )
  q05 = R.as__numeric(
    R.quantile(vec, probs: 0.05, names: false, type: 7))
  puts "5% quantile: #{q05 >> 0}"
end
```

```
## R class: Table
## 5% quantile: -0.00895
```

The Ledger's fallback path is similar: a small `data.frame` plus `R::Arrow.table_from(df)` when IPC is unavailable.

5.2 Example: IPC file handoff

Preferred for larger same-machine panels. Ruby writes the file; R opens by path; Ruby can unlink after R has the table. The reverse direction (`R::Arrow.write_ipc` then `Galaaz::ArrowIpc.read`) is how the Ledger can pull density coordinates back for Plotly.

```
ipc_ok = Galaaz::ArrowIpc.available? &&
  (R::Support.eval(
    "requireNamespace('arrow', quietly=TRUE)" == true)
unless ipc_ok
  puts '(Skip: need Arrow IPC backend + R arrow.)'
else
  returns = [-0.01, 0.02, 0.005, -0.003, -0.008]
  path = Galaaz::ArrowIpc.write(
    'daily_return' => returns.map(&:to_f))
  begin
    tbl = R::Arrow.open_ipc(path)
    puts "R class: #{tbl.rclass}"
    puts "IPC: #{File.basename(path)}"
    vec = R.as__numeric(
      R.dplyr__collect(tbl)[['daily_return']]
    )
    q05 = R.as__numeric(
      R.quantile(vec, probs: 0.05, names: false, type: 7))
    puts "5% quantile: #{q05 >> 0}"
  ensure
    Galaaz::ArrowIpc.release(path)
  end
end
```

```
## R class: Table
```



```
## IPC: galaaz_ipc_34399_93fe645f9fa9172f.arrow  
## 5% quantile: -0.0096
```

Then the product rule is **Remote Control**: keep the heavy panel in R; unbox **KPIs and chart coordinates** back to Ruby for SQLite and the browser.

```
SQLite -> portfolio return series (Ruby)  
        -> Arrow IPC file (path on the bridge)  
        -> GNU R (quantile / density / GBM)  
        -> JSON-ish payloads on SQLite  
        -> Turbo Stream -> Stimulus -> Plotly.js
```

6 The page with the charts

After a Local R stress test, the portfolio show page replaces the results panel in place. Captures below are from a live run of the demo (synthetic seed, Local engines, Plotly in the browser—not ggplot2 SVG).

6.1 Side-by-side KPIs

Engine A is **1-day historical** VaR. Engine B is **30-day forward** Monte Carlo. The UI says so on purpose: compare directionally, not as identical metrics.

6.2 Density and tail risk

Historical kernel density from Engine A, Monte Carlo terminal-return density from Engine B, with VaR markers as vertical lines.

6.3 Thirty-day Monte Carlo cone

Sample paths, a 5–95% band, median, and baseline—forward uncertainty as a picture, not only a scalar.

6.4 Rolling historical VaR and breaches

Trailing returns against a rolling 95% VaR line, with breach markers when returns punch through the threshold.

7 Why this is a production-shaped story

- **One desk.** Rails conventions for the app; CRAN for the science.
- **Jobs, not request-thread math.** Stress work runs in `StressTestJob`; the UI waits on Turbo/Cable.
- **Process isolation.** GNU R is a separate process (NewBridge). A bad R call need not take down Puma.
- **Optional dual-version R.** The same job can target Docker images (e.g. R 3.6.3 || 4.3.3) when you want version isolation; Local alone still tells the Arrow + DSL story.




Figure 1: Ledger results panel: Local R runtime banner, KPI table, and three Plotly charts.




Figure 2: Density and tail risk (Plotly): historical vs Monte Carlo terminal returns.




Figure 3: 30-day Monte Carlo cone (Plotly): paths, median, and 5-95% band.




Figure 4: Rolling historical VaR and breaches (Plotly).

- **OnRails family.** Ruby on Rails -> Omarchy / LinuxOnRails -> **R-on-Rails**—same “one person can ship the whole product” energy, applied to people whose science already lives in R.

8 Try it

```
git clone https://github.com/rbotafogo/r_on_rails_ledger.git
cd r_on_rails_ledger
bundle install
bin/rails db:prepare
SEED_PROFILE=fast bin/rails db:seed
bin/dev
# open http://localhost:3000 -> portfolio -> Run stress test
```

With a Galaaz checkout nearby, `galaaz add ledger` (Omarchy / CLI helper) can clone and boot the same demo. In-app docs live at `/docs` (architecture, Ruby DSL excerpts, runbook).

For the bridge itself—JRuby/CRuby, Arrow handoffs, `R::Job`—see the companion post **Galaaz 2.0** in this blog series.

9 Honest limits

- Seed prices are **synthetic**, not broker data.
- VaR / GBM here are a **demo risk stack**, not a bank production engine (no full PerformanceAnalytics suite in the live path).
- Charts are **Plotly.js** fed by JSON coordinates—not ggplot embedded as SVG in this build.
- The Ledger’s Arrow path is an **IPC file / mmap handoff**, not shared Ruby/R heap memory (zero-copy is still future work).
- Historical and Monte Carlo **horizons differ**; the KPI footnotes exist so the demo does not oversell a single number.

The point of honesty is the same as the point of the architecture: Rails can be the product, R can stay R, and you can show the charts without pretending the hard parts vanished.